

Where Does Bias Hide in Neural Retrieval?

Token-Level Attribution of Identity-Induced Score Sensitivity in Late-Interaction Models

Anonymous

Abstract

Neural retrieval models can exhibit score sensitivity when identity markers are swapped in documents, yet existing audits report only aggregate score differences without explaining where these shifts originate. We leverage ColBERT’s decomposable late-interaction scoring to perform a token-level audit of score sensitivity in late-interaction dense retrieval. We introduce **Token Contribution Disparity (TCD)**, a metric that exactly decomposes counterfactual retrieval-score shifts down to individual query-token contributions. Across **55,440 controlled tests**, supplemented by naturalistic-template and MS MARCO validations, we find that: (1) identity-induced score shifts disproportionately concentrate in function words rather than content words, with function words absorbing $1.43\times$ to $2.08\times$ more score change; (2) this pattern is robust across templates, query syntax, and real-world web text; (3) gender mismatch is the strongest predictor of absolute score sensitivity ($1.26\times$ higher than same-gender pairs), whereas cross-race pairing alone does not independently increase absolute sensitivity; and (4) subword fragmentation from tokenization yields a modest but consistent amplification. We show that while these score sensitivities are absent in BM25, we observe non-zero aggregate score sensitivity across several neural retrieval architectures, while exact token-level decomposition remains specific to ColBERT. Finally, synthetic candidate-pool ranking simulations demonstrate that these score shifts can produce rank displacements and top- k membership changes.

1 Introduction

Modern search systems increasingly rely on neural retrieval models that encode queries and documents as dense vector representations. In high-stakes domains such as hiring platforms and professional search, these models score candidate profiles against queries like “Who is a qualified doctor?” by

computing learned similarity between vector representations (Singh and Joachims, 2018). When the model processes documents describing people, identity cues such as names and pronouns can influence the computed relevance scores, potentially leading to unfair outcomes.

Prior work has documented that neural rankers exhibit measurable identity-related score sensitivity (Rekabsaz et al., 2021; Ziems et al., 2024; Manchanda and Shivaswamy, 2025). However, these studies report *aggregate* score differences: they tell us *that* a model is sensitive, but not *how*. We address this gap by exploiting ColBERT’s late-interaction MaxSim scoring, where the total relevance score decomposes into a sum of query-token contributions. This allows us to define Token Contribution Disparity (TCD), an exact token-level diagnostic that attributes counterfactual score shifts to individual query tokens. Unlike post-hoc attribution methods, TCD follows directly from the model’s scoring function.

We organize our study around four key research questions:

RQ1 Do identity swaps induce systematic token-level score shifts in ColBERT? (§4.1)

RQ2 Which query tokens carry identity-induced score shifts—function words or content words? (§4.2)

RQ3 What modulates the name-swap effect—race-associated pairing, gender-associated mismatch, or tokenization? (§4.3, §4.4)

RQ4 Is aggregate identity sensitivity specific to ColBERT, or does it also appear across other neural retrieval architectures? (§4.5)

Our main findings are: (1) function words absorb $1.43\text{--}2.08\times$ more identity-induced score change than content words ($d = 0.36\text{--}0.94$), a pattern observed across 29/33 professions; (2) non-identity

perturbation controls show that this score sensitivity carrier behavior extends to semantically weak, highly fragmented tokens; (3) gender mismatch is the most stable observed predictor of score sensitivity magnitude ($1.26\times$ in the partition analysis), whereas cross-race pairing is not a reliable predictor of absolute sensitivity after controls; and (4) synthetic candidate-pool ranking simulations show non-trivial rank and top- k changes under identity swaps.

2 Related Work

Fairness in information retrieval. Singh and Joachims (2018) formalized exposure fairness in rankings. Rekabsaz et al. (2021) measured gender bias in BERT-based rankers and proposed adversarial mitigation at the aggregate level. Krieg et al. (2023) created the Grep-BiasIR dataset for gender bias evaluation. Ziems et al. (2024) proposed the PAIR framework for auditing indexical bias. Blodgett et al. (2020) argued that “bias” in NLP must be carefully specified with respect to who is harmed. To our knowledge, our work is the first to exploit ColBERT’s MaxSim decomposition for token-level score-sensitivity auditing, differing from prior retrieval fairness studies that focus on aggregate measurement.

Counterfactual auditing and names as proxies. Bertrand and Mullainathan (2004) established the name-swap audit methodology for labor market discrimination. Sweeney (2013) demonstrated similar effects in online advertising. Caliskan et al. (2017) showed that static embeddings absorb human-like biases. De-Arteaga et al. (2019) demonstrated stereotyped representations in contextualized models applied to occupation prediction. Manchanda and Shivaswamy (2025) found that names distort embedding similarity and proposed anonymization. We adopt the name-swap framework but go beyond aggregate scores to analyse *per-token* effects and decompose the confounded name signal.

Tokenization and fairness. Petrov et al. (2023) showed that BPE tokenizers introduce cost and latency unfairness across languages. Ovalle et al. (2024) demonstrated that BPE over-fragmentation of neopronouns causes LLM misgendering. Bostrom and Durrett (2020) showed that tokenization choices affect downstream model behavior. Our rarity amplification finding extends this line of work to the retrieval scoring domain.

Late-interaction retrieval. ColBERT (Khattab and Zaharia, 2020) and ColBERTv2 (Santhanam et al., 2022) preserve per-token representations and compute scores via MaxSim. This decomposability makes them uniquely suited for token-level attribution—a property that single-vector models like DPR (Karpukhin et al., 2020) and Contriever (Goldfarb-Tarrant et al., 2024) lack. We leverage this property both as the subject and the instrument of our score-sensitivity audit.

Interpretability and token-level attribution.

Post-hoc token-level explanation methods—including LIME (Ribeiro et al., 2016) and integrated gradients (Sundararajan et al., 2017)—can provide approximate per-token importance scores for any model. Jain and Wallace (2019) cautioned that raw attention weights are not reliable explanations. SPLADE (Formal et al., 2021) computes sparse per-term weights that provide a form of token-level scoring. We explore attention *rollout* (Abnar and Zuidema, 2020) to test whether differential attention routing explains the function-word sensitivity pattern. Unlike post-hoc approximations, our TCD metric provides an *exact* decomposition of ColBERT’s MaxSim score by construction.

3 Methodology

3.1 Task and Model

We study counterfactual score sensitivity in candidate retrieval settings: ranking person-describing documents—for example, resumes, candidate profiles, or expert biographies—in response to professional queries (e.g., “Who is a qualified doctor?”).

We use ColBERTv2 (Santhanam et al., 2022), which encodes queries and documents into per-token L2-normalized embeddings using a shared BERT encoder. Relevance is computed via **MaxSim**:

$$S(Q, D) = \sum_{i=1}^{|Q|} \max_{j \in [1, |D|]} \mathbf{q}_i^\top \mathbf{d}_j \quad (1)$$

where $Q = (\mathbf{q}_1, \dots, \mathbf{q}_{|Q|})$ is the sequence of query token embeddings, $D = (\mathbf{d}_1, \dots, \mathbf{d}_{|D|})$ is the document token sequence, and $\mathbf{q}_i^\top \mathbf{d}_j$ is the cosine similarity between the i -th query token and j -th document token.

Intuitively, each query token independently searches for its single best-matching document token, and the total score is the sum of these best




Figure 1: **Overview of the TCD auditing pipeline.** *Step 1:* We construct counterfactual document pairs that are identical except for an identity name (e.g., Emily $\rightarrow$ Lakisha). *Step 2:* ColBERT encodes each query token and document token independently, then computes per-token MaxSim scores—each query token finds its single best-matching document token. *Step 3:* Absolute TCD measures sensitivity magnitude: $|\text{MaxSim}(D_A)_i - \text{MaxSim}(D_B)_i|$ (where $\text{MaxSim}(D)_i = \max_j \cos(q_i, d_j)$). We classify query tokens as **function words** (who, is, an, ?) or **content words** (experienced, doctor) and aggregate. The ratio of Func-TCD to Cont-TCD is $1.43\times$ in the controlled experiment and $2.08\times$ in the naturalistic-template validation.

matches. This *additive* structure is what enables per-token attribution: the contribution of query token q_i to the total score is simply $c_i = \max_j \mathbf{q}_i^\top \mathbf{d}_j$.

A key property is that document encoding is independent of other documents: $S(Q, D)$ does not depend on other candidates. Our pairwise comparisons are therefore mathematically identical to the scores assigned in a full retrieval index.

3.2 Counterfactual Design

To isolate the effect of identity markers on retrieval scores, we adopt a counterfactual design that holds document content constant while varying the identity cue. A **counterfactual pair** (D_A, D_B) consists of two documents identical except for an identity marker (note that the name also carries frequency, tokenization, and phonological information, which we decompose in §4.4):

D_A : “Emily has over ten years of experience in this field.”

D_B : “Lakisha has over ten years of experience in this field.”

The per-token contribution of query token q_i is $c_i^A = \max_j \mathbf{q}_i^\top \mathbf{d}_j^A$ for document D_A , and analogously c_i^B for D_B . In plain terms, c_i is the score contribution of query token q_i —the highest similarity it achieves with any document token.

3.3 Token Contribution Disparity (TCD)

Aggregate score gaps tell us that a model treats two documents differently, but not where the difference enters in the query-token decomposed scores. TCD answers this by accounting for the score gap at the level of individual query-token contributions.

For a query q with tokens $q_1, \dots, q_n$ and a counterfactual pair (D_A, D_B) , ColBERT scores decompose as:

$$S(q, D) = \sum_{i=1}^n \max_j \cos(q_i, d_j) \quad (2)$$

The **Token Contribution Disparity** for query token q_i is:

$$\text{TCD}_i = \max_j \cos(q_i, d_j^A) - \max_j \cos(q_i, d_j^B) \quad (3)$$

We define signed TCD as the per-token contribution difference. Unless otherwise stated, our analyses use absolute TCD, $|\text{TCD}_i|$, to measure sensitivity magnitude. TCD is an *exact accounting* of the total score difference: $S(q, D_A) - S(q, D_B) = \sum_i \text{TCD}_i$. We distinguish this exact accounting from causal attribution or social harm measurement; we pair TCD with perturbation controls, signed group analyses, and ranking simulations to address each dimension. We aggregate by token type to form Func-TCD and Cont-TCD.

To quantify overall score sensitivity, we define **Score Sensitivity** (SS), the relative absolute score difference:

$$SS = \frac{|S(q, D_A) - S(q, D_B)|}{\frac{1}{2}(S(q, D_A) + S(q, D_B))} \quad (4)$$

3.4 Name Pool

We use **65 names** drawn from the Bertrand and Mullainathan (2004) audit set and the Rosenman et al. Harvard Dataverse ethnicity database, covering four US-associated racial categories: White (10 female, 9 male), Black (9 female, 12 male), Hispanic (6 female, 8 male), and Asian (5 female, 6 male). These form **120 stratified counterfactual name pairs**, balanced across cross-race same-gender, cross-gender same-race, within-group, and matched-rarity categories. Each name is annotated with race, gender, and subword token count (computed using ColBERTv2’s WordPiece tokenizer; see Appendix A). Names tokenized into a single subword are classified as *common*; those requiring >1 subword are *rare*. We treat these labels as name-associated audit variables, not deterministic demographic identities.

3.5 Experimental Setup

Queries. Each query follows the template “Who is an experienced {profession}?”, instantiated with 33 professions spanning medicine (doctor, surgeon, nurse, dentist, therapist), STEM (software engineer, data scientist, civil engineer, biologist, researcher), education (teacher, professor, librarian), law/business (lawyer, accountant, financial analyst, consultant, CEO, manager, director), trades (electrician, plumber, mechanic, construction worker, truck driver, welder), and service roles (secretary, social worker, receptionist, janitor, chef, pilot, firefighter). This range was chosen to cover high-prestige, gender-stereotyped, and blue-collar domains, ensuring the analysis is not limited to any single sector.

We use **14 document templates** that simulate professional profiles or biographical summaries, varying in three dimensions: (1) the *syntactic position* of the identity marker (sentence-initial vs. mid-sentence vs. after a preposition), (2) the *density of function words* surrounding the marker, and (3) the *register* (formal bio vs. resume header vs. peer endorsement). For example, template T1 (“{NAME} has over ten years of experience...”) places the name sentence-initially with few surrounding function words, while T7 (“Resume of {NAME}, an

experienced...”) places it after a preposition and determiner. All 14 templates appear in Appendix C.

Scale. Crossing 33 queries $\times$ 120 name pairs $\times$ 14 templates yields **55,440 counterfactual tests**. We supplement the main experiment with: (i) a control validation comparing identity swaps to non-identity substitutions (“ten” $\rightarrow$ “fifteen”); (ii) a BM25 baseline ($n=1,155$); (iii) cross-architecture comparison on the same counterfactual pairs using DPR, SPLADE, Contriever, and a cross-encoder ($n=1,155$ each); (iv) a name-masking sanity check; (v) a naturalistic-template validation with 200 diverse generated passages ($n=8,000$; §4.8); and (vi) an ecological validation on naturally occurring passages from the MS MARCO benchmark ($n=5,000$; §4.9). All code and data are available anonymously.¹

Statistical treatment. Because tests reuse names, professions, and document templates, the 55,440 observations should not be interpreted as independent draws. We therefore emphasise effect sizes, ratios, mixed-effects robustness, and cluster-aware checks over raw p -values. Non-parametric paired tests are reported as descriptive evidence for the paired counterfactual design; confound regressions include profession and template controls with clustered standard errors, and we report whether effects remain stable under these controls.

4 Results

4.1 Control Validation (RQ1)

Identity swaps produce substantially higher FuncTCD than a simple non-identity substitution control (“ten” $\rightarrow$ “fifteen”; Cohen’s $d = 1.86$, $p = 7 \times 10^{-6}$).²

Non-identity perturbation control. To test whether function-word amplification is identity-specific, we compare five perturbation types: identity name swaps, numeric word swaps, adjective swaps, institution swaps, and rare nonsense words (not real names but heavily subword-fragmented). Identity swaps ($F/C = 1.81\times$, $d = 0.74$) and rare nonsense words ($1.67\times$, $d = 0.68$) both show strong amplification, while number ($1.00\times$), adjective ($1.05\times$), and institution swaps ($0.71\times$) show

¹<https://anonymous.4open.science/r/colbert-bias-audit>

²Extremely small p -values throughout this paper reflect repeated paired measurements and large sample sizes; effect sizes and cluster-aware robustness checks are the more informative evidence.

none. The effect is therefore not unique to identity labels, but is consistent with a broader sensitivity to semantically weak or heavily fragmented subject-position tokens. Identity names instantiate this sensitivity in a socially consequential setting because they also carry race-, gender-, and frequency-associated signals.

4.2 Function-Word Concentration of Identity-Induced Score Shifts (RQ2)

Across 55,440 counterfactual tests (33 professions $\times$ 120 name pairs $\times$ 14 templates), function words consistently exhibit higher TCD than content words (Figure 2):

- Func-TCD = 0.018 ± 0.013 ; Cont-TCD = 0.013 ± 0.011
- **Ratio = $1.43\times$** (cluster-robust 95% CI: [1.35, 1.51]; Cohen’s $d = 0.36$)
- **29/33 professions (88%)** exhibit the pattern (Figure 3)

Cluster bootstrap checks confirm that the ratio remains above 1.0 when resampling whole name pairs (95% CI: 1.35–1.51), professions (1.30–1.57), templates (1.30–1.56), or name-pair $\times$ profession cells (1.40–1.47). A crossed mixed-effects model with random intercepts for name pair, profession, and template estimates a positive function-word effect ($\beta = 0.00544$, 95% CI [0.00534, 0.00555]); a paired-difference random-effects model gives a wider but still positive interval ([0.00447, 0.00642]). Thus the main pattern is not an artifact of treating all 55,440 rows as independent. Additional token-grouping robustness checks in Appendix E (Table 7) show that the result is stable under punctuation removal, stopword-list, POS-based, and closed-class grouping definitions.

The four exceptions—janitor ($r=0.84$), electrician ($r=0.87$), plumber ($r=0.91$), and social worker ($r=0.99$)—are cases where the content word itself appears more sensitive to the identity swap, reducing or reversing the function-word dominance.

Template modulation. The ratio varies with document structure (Table 1). Across 14 templates, the ratio ranges from $0.96\times$ (T12) to $1.84\times$ (T11). When the identity marker appears in a context with high function-word density (e.g., “Certified and licensed, {NAME}...”), the ratio is highest. 13 of 14 templates produce ratios above 1.0.




Figure 2: Distribution of TCD values for function words vs. content words across 55,440 counterfactual tests. Function words show consistently higher score sensitivity ($1.43\times$, cluster-robust 95% CI [1.35, 1.51]; $d = 0.36$).

	Template (abbreviated)	Ratio
T1	{NAME} has over ten years...	$1.41\times$
T2	{NAME} is a dedicated...	$1.67\times$
T3	{NAME} brings extensive...	$1.56\times$
T4	The candidate {NAME}...	$1.41\times$
T5	As a leading expert, {NAME}...	$1.32\times$
T6	{NAME} graduated from...	$1.53\times$
T7	Resume of {NAME}...	$1.26\times$
T8	{NAME} completed a fellowship...	$1.10\times$
T9	A highly rated... {NAME}...	$1.31\times$
T10	With dual expertise... {NAME}...	$1.60\times$
T11	Certified and licensed, {NAME}...	$1.84\times$
T12	{NAME}'s career spans...	$0.96\times$
T13	Colleagues describe {NAME}...	$1.45\times$
T14	Having worked internationally, {NAME}...	$1.82\times$
Combined (14 templates)		$1.43\times$

Table 1: Func/Cont TCD ratio by document template. 13/14 templates produce ratios above 1.0. The ratio ranges from $0.96\times$ to $1.84\times$, modulated by the local function-word density surrounding the identity marker.

This modulation suggests that the sensitivity pathway is mediated by the *local context* of the identity marker, not merely its presence. Function words near the name are more strongly affected, a pattern consistent with contextual representations carrying name-derived information into nearby tokens.

4.3 Rarity Amplification (RQ3a)

We stratify name pairs by subword token count: CC (both common, 1 token each), CR (mixed), and RR




Figure 3: Func/Cont TCD ratio for each of 33 professions. Coral bars indicate Func > Cont; purple bars indicate the four exceptions. 29 out of 33 professions show higher function-word TCD.

(both rare, >1 token each). Pairs where both names are rare show a modest but consistent increase in score sensitivity:

- CC: SS = 0.028; RR: SS = 0.031
- **RR/CC ratio = 1.11×** (Cohen’s $d = 0.13$)

The amplification occurs in the *absolute rarity* condition—when both names undergo subword fragmentation—rather than in the token *gap* between names. Since minority-group names are disproportionately fragmented by WordPiece (Petrov et al., 2023), this suggests a potential infrastructure-level fairness concern associated with tokenization fragmentation (Figure 4).

4.4 Confound Decomposition (RQ3b)

To disentangle race, gender, and tokenization effects, we employ nested OLS regression with clustered standard errors (clustered by profession) alongside the partition analysis. Table 2 shows the simple partition results; the regression provides partial coefficients that control for confound correlations.

Confound	n	Mean SS	Comparison
Same-gender	35,112	0.027	baseline
Cross-gender	20,328	0.035	1.26×
Within-race	20,328	0.031	baseline
Cross-race	35,112	0.030	0.97×

Table 2: Confound decomposition (partition analysis). Gender mismatch produces the largest score sensitivity increase (1.26×, $d = 0.28$); cross-race pairing alone does not amplify absolute sensitivity.




Figure 4: Rarity amplification by name-pair category. Rare–rare (RR) pairs show a modest but monotonic increase in score sensitivity over common–common (CC) pairs. Points show means; whiskers show 95% bootstrap CIs.

Regression analysis. Nested OLS regression with profession and template controls confirms gender mismatch as the most stable partial predictor of Score Sensitivity ($\beta = 0.010$, $p < 10^{-14}$). Cross-race pairing is not significant after these controls ($\beta = 0.001$, $p = 0.40$), while name-frequency disparity is smaller but significant ($\beta = 0.003$, $p = 0.027$). The full controlled model explains roughly one fifth of observed SS variation, indicating that profession, template, and name-level features jointly shape the effect.

Three findings emerge (Figure 5): (1) **Gender mismatch** shows the strongest observed association with score sensitivity (1.26× higher SS, $d = 0.28$). (2) **Cross-race pairing** does not independently amplify absolute score sensitivity (ratio = 0.97× in the partition analysis; n.s. after controls). (3) **Subword rarity** adds modest amplification (RR/CC = 1.11×, $d = 0.13$).

Gender also modulates the function-word ratio: cross-gender pairs exhibit a Func/Cont ratio of 1.88× compared to 1.35× for same-gender pairs. This suggests that gender incongruity between the

Model	Type	SS	Token?
BM25	Lexical	0.000	—
ColBERTv2	Late-inter.	0.036	Yes
Cross-encoder	Joint	0.050	No
Contriever	Bi-encoder	0.054	No
DPR	Bi-encoder	0.100	No
SPLADE	Sparse	0.123	No

Table 3: Identity sensitivity (SS) across retrieval architectures ($n=1,155$ per model). BM25 shows zero sensitivity; all neural models show non-zero sensitivity under their own scoring functions. Because score scales differ across architectures, the values should not be read as a calibrated fairness ranking. Only ColBERT supports exact token-level TCD decomposition.




Figure 5: **Partition analysis for gender- and race-associated pairings.** Bars show mean Score Sensitivity; error bars show 95% bootstrap CIs. Cross-gender pairs show higher Score Sensitivity than same-gender pairs, whereas cross-race pairs do not exceed within-race pairs.

query context and the identity marker intensifies the function-word sensitivity pattern.

4.5 Cross-Architecture Validation (RQ4)

To test whether identity sensitivity is specific to ColBERT, we run the same counterfactual pairs through BM25 and five neural retrieval architectures ($n=1,155$ each; Table 3).

BM25 produces $SS = 0.000$ across all 1,155 tests, confirming that lexical matching alone is insensitive to identity swaps when all non-name terms are unchanged. All five neural models show non-zero identity sensitivity under their own scoring functions. Because these architectures use different score scales, we interpret Table 3 as evidence that aggregate identity sensitivity is not unique to ColBERT, rather than as a calibrated ranking of model fairness. Only ColBERT enables the exact token-level decomposition that reveals *where* the score difference enters (see Figure 8 in the Appendix).

4.6 Mechanistic Diagnostics

Name masking (sanity check). Replacing all name tokens in the *documents* with [MASK] makes the two documents in each counterfactual pair identical, and TCD drops to exactly zero. While this outcome is expected by construction, it verifies that the implementation attributes score changes to the name-bearing counterfactual difference rather than to unrelated bookkeeping or positional artifacts.

MaxSim argmax alignment. To test whether function-word TCD is simply caused by query tokens switching to different best-matching document tokens, we trace the MaxSim argmax target for each query token across 11,088 counterfactual document pairs ($24 \text{ name pairs} \times 33 \text{ professions} \times 14 \text{ templates}$; 70,560 query-token observations; Appendix E). Function words still show higher mean TCD than content words ($1.35\times$), but they do *not* switch argmax targets more often: their argmax-change rate is slightly lower than content words (0.59 vs. 0.61). Nor is the effect explained by direct matching to the name span or its local context: function words hit the name ± 2 -token window less often than content words (0.35 vs. 0.43). This negative result, visualised in Appendix Figure 7, rules out a simple “function words jump to the name token” explanation and instead is consistent with contextual representational shifts before the final MaxSim operation, but does not establish a complete causal pathway.

Attention rollout. We apply attention rollout (Abnar and Zuidema, 2020) to test whether name tokens route more attention to function-word positions than to content-word positions. The result is *negative*: the attention ratio is 1.00 ($p = 0.065$). Layer-by-layer embedding probing suggests a more distributed representational pathway: at intermediate layers (L1–L9), content-word embeddings shift more than function words (ratio 0.41–0.52 $\times$), but at the final layer (L12) function words converge to a 1.07 $\times$ ratio. We therefore treat the mechanism as partially characterised rather than fully established.

Query ablation probe. To test whether function words merely appear as post-hoc attribution carriers or also affect aggregate score sensitivity, we compare the original function-heavy query (“Who is an experienced {profession}?”) with a content-only ablated query (“experienced {profession}”) across a representative subset of 700 tests. The original query yields slightly higher Score Sensi-

tivity (SS) than the ablated query (mean SS 0.0361 vs. 0.0344; ratio = $1.051\times$; paired Wilcoxon $p = 1.31 \times 10^{-2}$; Cohen’s $d = 0.088$). This small but statistically reliable reduction suggests that function-word context contributes to overall sensitivity, although the effect size is modest and the ablation changes query length and syntax as well as function-word content.

4.7 Ranking Impact

Identity-induced score differences translate into ranking consequences. Across 55,440 counterfactual tests, **77.7%** exhibit a non-trivial score difference (>0.01). In synthetic candidate-pool ranking simulations designed to approximate realistic qualification variance, identity swaps change at least one rank position in 36.6% (95% CI: [33.7%, 39.7%]), 55.0% ([52.6%, 57.1%]), and 75.1% ([73.7%, 76.4%]) of cases for pool sizes of 10, 20, and 50 documents, respectively. Mean absolute displacement grows with pool size (0.58 [0.52, 0.64], 1.21 [1.12, 1.29], and 3.07 [2.92, 3.21] positions), and top- k membership changes are most visible in the Top-1 (8.5%, 21.0%, 41.0%) and Top-3 (9.8%, 7.8%, 18.3%) sets (Table 10). Appendix Table 11 further shows that flips heavily concentrate in small-margin adjacent pairs; for pool size 50, 83.9% of adjacent pairs have a score margin of < 0.05 , where observed rank flip rates reach 25.6%–31.0%.

Signed directional analysis. Signed score gaps are necessary for distinguishing identity sensitivity from group disadvantage. In the expanded name pool, the direction of score gaps is substantially name-pair dependent: cross-race pairing does not increase absolute score sensitivity after controls, even though individual name pairs can show large signed gaps. We therefore use the signed analysis as a diagnostic for deployment-specific exposure risk rather than as evidence for a global group ordering.

Gender direction. Male-associated names score higher than female-associated names in 62.3% of cross-gender tests (mean gap $+0.045$). This directional pattern is consistent with the stronger absolute sensitivity observed for cross-gender pairs, but should be interpreted as a property of this name pool and query set rather than a universal statement about gender.




Figure 6: Function-word sensitivity in controlled templates vs. naturalistic-template passages. The effect strengthens on longer diverse passages ($2.08\times$ vs. $1.43\times$, $d=0.94$ vs. 0.36), suggesting that the controlled setup is not the only source of the pattern.

4.8 Naturalistic-Template Validation

To test whether the function-word sensitivity pattern generalizes beyond our 14 short templates, we apply the identical TCD pipeline to **200 diverse professional passages** generated from 16 structurally distinct templates spanning 14 professions and 13 knowledge domains (see Appendix D for examples). For each passage we perform the same 5-pair $\times$ 8-query counterfactual analysis, yielding **8,000 additional tests**. The function-word pattern strengthens on these longer naturalistic-template passages (Figure 6): the Func/Cont TCD ratio is **$2.08\times$** (Cohen’s $d = 0.94$), compared to $1.43\times$ on the controlled templates. All 8/8 tested query professions show a ratio above 1.0. Removing punctuation from the function-word category does not weaken the pattern ($F/C = 1.58\times$ without “?”), confirming the result is not driven by punctuation tokens.

4.9 MS MARCO Benchmark Validation

To verify whether the function-word sensitivity pattern persists in naturally occurring text, we perform an ecological validation using the standard **MS MARCO passage ranking collection** (Nguyen et al., 2016). We run a SpaCy NER scan over the first 10,000 passages in the collection, identifying passages containing person name entities. For the first 100 naturally occurring passages containing a person name, we perform the counterfactual name-swap audit using our pool of 10

names and 5 professional queries, yielding **5,000 counterfactual tests**.

We find that the function-word sensitivity pattern holds robustly on naturally occurring MS MARCO text. The mean Func-TCD is **0.0162**, whereas the mean Cont-TCD is **0.0098**, yielding a Func/Cont TCD ratio of **1.65** $\times$. Function-word amplification occurs in **70.1%** of all comparisons, and the overall mean Score Sensitivity is **0.0276**. This result demonstrates that our findings are not an artifact of template-generated documents, but rather represent a robust representational pattern of contextual late-interaction models processing real-world web text.

5 Discussion

Function-word sensitivity. Non-identity perturbation control (§4.1) shows function-word amplification extends to rare nonsense words ($1.67\times$) but not semantically grounded substitutions, consistent with sensitivity to semantically weak or fragmented tokens. Identity names instantiate this sensitivity in a socially consequential setting because name tokens carry demographic associations as well as frequency and tokenization structure. The naturalistic-template validation ($2.08\times$) and template modulation ($0.96\times$ – $1.84\times$) support generalizability beyond a single short template.

Attention null and layer probing. Attention rollout did not confirm differential pathways ($p = 0.065$), while layer-by-layer probing and argmax alignment suggest that identity signal is redistributed across contextual representations before the final MaxSim computation. This points to a plausible representational pathway, but further head-level analysis is needed before making a stronger mechanistic claim.

Confound structure. Gender mismatch drives sensitivity magnitude ($1.26\times$), while cross-race pairing does not independently increase absolute score sensitivity after controls. Signed score gaps remain important for deployment-specific harm assessment, but in our expanded name pool they are strongly name-pair dependent rather than reducible to a single universal group ordering. Subword fragmentation (Petrov et al., 2023; Ovalle et al., 2024) modestly amplifies the effect, suggesting an infrastructure-level issue that content-level debiasing alone may miss.

ColBERT as diagnostic. Integrated gradients produces a concordant function/content ratio (IG: $1.56\times$, Spearman $\rho = 0.49$, $p < 10^{-44}$), providing an approximate post-hoc check that agrees with TCD’s exact MaxSim decomposition.

6 Conclusion

We have presented TCD, a token-level score-sensitivity diagnostic exploiting ColBERT’s decomposable MaxSim scoring. Across 55,440 controlled tests spanning 65 names from four US-associated racial categories, function words absorb $1.43\times$ more identity-induced score change than content words; this ratio increases to $2.08\times$ in the naturalistic-template validation. Gender mismatch is the most stable observed predictor of sensitivity magnitude ($1.26\times$), while cross-race pairing does not independently amplify absolute sensitivity after controls. Practical interventions include **document-side name masking** (which eliminates TCD but may sacrifice legitimate relevance), **equitable tokenizer design**, and **TCD monitoring** as a continuous sensitivity diagnostic.

Limitations

Synthetic core design. Our main experiment uses template-generated counterfactual pairs. While our naturalistic-template validation (§4.8) demonstrates that the function-word pattern generalizes to longer and more diverse generated passages with an even stronger effect ($2.08\times$), these passages are still constructed. We validated the generalizability of the function-word pattern on real naturally occurring web text by running a 5,000-test audit on the MS MARCO benchmark (§4.9), confirming a strong $1.65\times$ ratio. However, a larger scale evaluation on other organic corpora (e.g., actual job postings, resumes) with naturally embedded names would further strengthen external validity.

Name pool scope. Our 65-name pool covers four US-centric race-associated name categories (White, Black, Hispanic, Asian) drawn from established audit studies and the Rosenman ethnicity database. Names from other racial groups, geographies, cultural contexts, and non-binary gender identities are not represented. The stratified pairing strategy improves coverage but cannot fully eliminate confounds between race, gender, frequency, orthography, and tokenization.

Single diagnostic model. While we compare aggregate sensitivity across BM25 and five neural retrieval architectures, the token-level TCD analysis is specific to ColBERT’s MaxSim decomposition. We validate TCD against integrated gradients (Spearman $\rho = 0.49$, concordant F/C ratio), but extending token-level attribution to other architectures remains future work.

Rarity amplification. The rarity effect ($RR/CC = 1.11\times$, $d = 0.13$) is modest. A larger, more diverse name pool with controlled frequency, orthography, and tokenizer-fragmentation levels would strengthen this finding.

Attention null result. Our attention rollout analysis did not confirm the hypothesised differential attention pathway ($p = 0.065$). The mechanism by which function words absorb identity information may not be captured by head-averaged rollout; layer-specific or head-specific analysis could reveal more nuanced patterns.

Ethical Considerations

This work audits score sensitivity in a retrieval model using established name-swap methodology. The names used are drawn from published audit studies (Bertrand and Mullainathan, 2004) and the Rosenman et al. ethnicity database, and are *race-associated* and *gender-associated*, not deterministic demographic identifiers. Our racial and gender categories are simplifications of complex social constructs and are used solely as analytical variables for bias measurement, not as claims about identity.

The counterfactual documents are synthetic and do not describe real individuals. TCD is a diagnostic metric measuring score sensitivity, not a direct measure of social harm. Signed score gaps and ranking simulations are required to connect sensitivity to exposure differences; in our experiments, these directional effects can be substantial for individual name pairs but should not be interpreted as deterministic statements about demographic groups.

The counterfactual analyses should inform—but not replace—human-centered fairness assessments before deployment decisions. Document-side name masking eliminates TCD but may sacrifice legitimate name-based relevance, making universal masking inappropriate without context-specific evaluation.

We acknowledge that bias auditing itself can reinforce categorical thinking about identity. We follow the recommendation of Blodgett et al. (2020) to specify concretely who is harmed and how, rather than treating “bias” as an undifferentiated property of models.

Use of AI assistants. An AI coding assistant was used to assist with experiment scripting and LaTeX formatting. All scientific claims, experimental design, and analysis were conducted by the authors.

References

- Samira Abnar and Willem Zuidema. 2020. Quantifying attention flow in transformers. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4190–4197. Association for Computational Linguistics.
- Marianne Bertrand and Sendhil Mullainathan. 2004. [Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination](#). *American Economic Review*, 94(4):991–1013.
- Su Lin Blodgett, Solon Barocas, Hal Daumé, and Hanna Wallach. 2020. [Language \(technology\) is power: A critical survey of “bias” in NLP](#). In *Proceedings of ACL 2020*, pages 5454–5476. ACL.
- Kaj Bostrom and Greg Durrett. 2020. [Byte pair encoding is suboptimal for language model pretraining](#). In *Findings of EMNLP 2020*, pages 4617–4624. ACL.
- Aylin Caliskan, Joanna J. Bryson, and Arvind Narayanan. 2017. [Semantics derived automatically from language corpora contain human-like biases](#). *Science*, 356(6334):183–186.
- Maria De-Arteaga, Alexey Romanov, Hanna Wallach, Jennifer Chayes, Christian Borber, Alexandra Chouldechova, Sahin Geyik, Krishnamurthy Kenthapadi, and Adam Tauman Kalai. 2019. [Bias in bios: A case study of semantic representation bias in a high-stakes setting](#). In *Proceedings of FAccT 2019*, pages 120–128. ACM.
- Thibault Formal, Benjamin Piwowarski, and Stéphane Clinchant. 2021. [SPLADE: Sparse lexical and expansion model for first stage ranking](#). In *Proceedings of the 44th International ACM SIGIR Conference*, pages 2288–2292. ACM.
- Seraphina Goldfarb-Tarrant, Pedro Rodriguez, Jane Dwivedi-Yu, and Patrick Lewis. 2024. [MultiContrivers: Analysis of dense retrieval representations](#). In *Proceedings of BlackboxNLP 2024*, pages 87–100. ACL.
- Sarthak Jain and Byron C. Wallace. 2019. [Attention is not explanation](#). In *Proceedings of NAACL 2019*, pages 3543–3556. ACL.

- Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense passage retrieval for open-domain question answering. In *Proceedings of EMNLP 2020*, pages 6769–6781. ACL.
- Omar Khattab and Matei Zaharia. 2020. [ColBERT: Efficient and effective passage search via contextualized late interaction over BERT](#). In *Proceedings of the 43rd International ACM SIGIR Conference*, pages 39–48. ACM.
- Klara Krieg, Emilia Parada-Cabaleiro, Gertraud Medicus, Oleg Lesota, Markus Schedl, and Navid Rekasaz. 2023. Grep-BiasIR: A dataset for investigating gender representation bias in information retrieval results. In *Proceedings of CHIIR 2023*. ACM.
- Sahil Manchanda and Pannaga Shivaswamy. 2025. What is in a name? Mitigating name bias in text embedding similarity via anonymization. In *Findings of ACL 2025*. ACL.
- Tri Nguyen, Mir Rosenberg, Xia Song, Jianfeng Gao, Saurabh Tiwary, Rangan Majumder, and Li Deng. 2016. MS MARCO: A human generated machine reading comprehension dataset. *arXiv preprint arXiv:1611.09268*.
- Anaelia Ovalle, Palash Goyal, Jwala Dhamala, Matthew Jagielski, Kai-Wei Chang, and Aram Galstyan. 2024. Tokenization matters: Navigating data-scarce tokenization for gender inclusive language technologies. In *Findings of NAACL 2024*. ACL.
- Aleksandar Petrov, Emanuele La Malfa, Philip Torr, and Adel Bibi. 2023. Language model tokenizers introduce unfairness between languages. In *Advances in Neural Information Processing Systems*, volume 36. NeurIPS.
- Navid Rekasaz, Simone Kopeinik, and Markus Schedl. 2021. [Societal biases in retrieved contents: Measurement framework and adversarial mitigation of BERT rankers](#). In *Proceedings of the 44th International ACM SIGIR Conference*, pages 306–316. ACM.
- Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. 2016. [“Why Should I Trust You?”: Explaining the predictions of any classifier](#). In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1135–1144. ACM.
- Keshav Santhanam, Omar Khattab, Jon Saad-Falcon, Christopher Potts, and Matei Zaharia. 2022. [ColBERTv2: Effective and efficient retrieval via lightweight late interaction](#). In *Proceedings of NAACL 2022*, pages 3715–3734. ACL.
- Ashudeep Singh and Thorsten Joachims. 2018. Fairness of exposure in rankings. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 2219–2228. ACM.
- Mukund Sundararajan, Ankur Taly, and Qiqi Yan. 2017. Axiomatic attribution for deep networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 3319–3328. PMLR.
- Latanya Sweeney. 2013. [Discrimination in online ad delivery](#). *Communications of the ACM*, 56(5):44–54.
- Caleb Ziems, William Held, Jane Dwivedi-Yu, and Diyi Yang. 2024. [Measuring and addressing indexical bias in information retrieval](#). In *Findings of ACL 2024*, pages 12860–12877. ACL.

Race	Female	Male	Total	% Rare
Asian	5	6	11	64%
Black	9	12	21	43%
Hispanic	6	8	14	71%
White	10	9	19	32%
Total	30	35	65	49%

Table 4: Name pool summary (65 names, 4 races). % Rare = proportion requiring >1 WordPiece token. Minority-group names are disproportionately fragmented.

A Name Pool

Table 4 summarises the 65-name pool by race and gender. Names are drawn from [Bertrand and Mullainathan \(2004\)](#) and the Rosenman et al. Harvard Dataverse ethnicity database.

B Professions and Query Format

All 33 queries follow the format “Who is an experienced {profession}?”:

doctor, nurse, surgeon, therapist, dentist, software engineer, researcher, data scientist, civil engineer, biologist, teacher, professor, librarian, lawyer, accountant, financial analyst, consultant, CEO, manager, director, electrician, plumber, mechanic, construction worker, truck driver, welder, secretary, social worker, receptionist, janitor, chef, pilot, firefighter.

C Document Templates

The 14 document templates, with the identity marker position shown as {NAME}:

- T1: “{NAME} has over ten years of experience in this field and has worked with top organizations.”
- T2: “{NAME} is a dedicated professional with extensive expertise and a proven track record.”
- T3: “{NAME} brings extensive experience and a unique perspective to every project undertaken.”
- T4: “The candidate {NAME} holds advanced credentials and has received multiple awards.”
- T5: “As a leading expert, {NAME} has published widely and contributed to major initiatives.”
- T6: “{NAME} graduated from a top university and has been recognized for outstanding work.”
- T7: “Resume of {NAME}, an experienced professional with deep domain expertise.”

T8: “{NAME} completed a prestigious fellowship and mentored junior colleagues.”

T9: “A highly rated professional, {NAME} has served clients across multiple industries.”

T10: “With dual expertise in management and strategy, {NAME} leads cross-functional teams.”

T11: “Certified and licensed, {NAME} maintains the highest standards of professional practice.”

T12: “{NAME}’s career spans two decades of innovation and leadership in this domain.”

T13: “Colleagues describe {NAME} as reliable, insightful, and exceptionally detail-oriented.”

T14: “Having worked internationally, {NAME} brings a global perspective to every engagement.”

D Naturalistic-Template Validation Details

The naturalistic-template validation (§4.8) uses 200 passages generated from 16 structurally distinct templates (independent of the 14 main-experiment templates), spanning 14 professions (physician, surgeon, attorney, engineer, researcher, specialist, professor, therapist, analyst, consultant, practitioner, clinician, advisor, director) and 13 knowledge domains (cardiovascular medicine, corporate litigation, structural engineering, machine learning, public health, pediatric oncology, intellectual property, environmental policy, data science, neuroscience, family law, renewable energy, clinical psychology). Example passages:

“{NAME} is a board-certified physician with over 15 years of clinical experience. {NAME} specializes in cardiovascular medicine and has published more than 45 peer-reviewed articles in leading journals.”

“The report authored by {NAME}, a leading researcher in the field, highlights the urgent need for reform in public health. {NAME} presented the findings before Congress last month.”

Each passage undergoes the same 5-pair $\times$ 8-query counterfactual analysis as the main experiment, yielding 8,000 tests. The Func/Cont TCD ratio is $2.08 \times$ (Cohen’s $d = 0.94$), confirming that the function-word pattern is not an artifact of the main experiment’s template structure.

Cluster	n	Med.	95% CI
Name pair	120	1.43	[1.35, 1.51]
Profession	33	1.43	[1.30, 1.57]
Template	14	1.43	[1.30, 1.56]
Pair $\times$ prof.	3,960	1.43	[1.40, 1.47]

Table 5: Cluster bootstrap robustness for the main Func/Cont TCD ratio.

Model	Estimate (95% CI)
Long TCD: function coefficient	0.00544 [0.00534, 0.00555]
Paired diff.: Func-Cont intercept	0.00544 [0.00447, 0.00642]

Table 6: Mixed-effects robustness. Both models include crossed random intercepts for repeated name-pair, profession, and template factors; both effects are positive ($p < 10^{-27}$).

E Robustness and Mechanism Details

Table 5 reports cluster bootstrap confidence intervals for the main Func/Cont TCD ratio. All intervals remain above 1.0 when resampling whole clusters rather than individual rows. Table 6 reports mixed-effects robustness checks that model repeated name-pair, profession, and template factors directly.

Token Grouping Robustness. To evaluate the robustness of our token classification scheme, we analyze alternative grouping criteria. Because of the structured and short nature of our query template (“Who is an experienced {profession}?”), standard NLP stopword lists (e.g., NLTK, scikit-learn), POS-tag based divisions, closed-class groupings, and punctuation-removal filters all yield highly consistent results. Specifically, excluding the question mark (?) results in a Func/Cont ratio of $1.58\times$ across the main experiment (since only the function words “who”, “is”, “an” are kept), while POS-tag or stopword groupings including the punctuation yield the baseline $1.43\times$ ratio. Across all classification schemes, function-word score sensitivity remains consistently and significantly higher than content-word sensitivity, confirming that our findings are robust to reasonable grouping definitions. Table 7 summarizes these grouping schemes and ratios.

Table 8 reports the MaxSim argmax-alignment check. Function words have higher TCD without higher argmax-switch or name-context-hit rates, ruling out a simple direct-matching explanation.

Grouping Scheme	Function Tokens	Content Tokens	Ratio
Main (Standard)	who, is, an, ?	experienced, {prof}	$1.43\times$
No punctuation	who, is, an	experienced, {prof}	$1.58\times$
Stopword list	who, is, an	experienced, {prof}	$1.58\times$
POS-based	who, is, an, ?	experienced, {prof}	$1.43\times$
Closed-class	who, is, an	experienced, {prof}	$1.58\times$

Table 7: Func/Cont TCD ratio across different query token grouping schemes. Due to the short query structure, all standard NLP classifications converge to either the $1.43\times$ or $1.58\times$ ratio, demonstrating robust amplification in all settings.

Metric	Function	Content
Mean TCD	0.0175	0.0129
Argmax-change rate	0.59	0.61
Name-span hit rate	0.001	0.194
Name ± 2 context hit rate	0.35	0.43
Mean TCD if argmax changed	0.0181	0.0152
Mean TCD if argmax stable	0.0166	0.0093

Table 8: MaxSim argmax-alignment results over 11,088 document pairs and 70,560 query-token observations.




Figure 7: Aggregate MaxSim argmax-alignment diagnostics. Function words show higher TCD but not higher argmax-switch or name-context-hit rates, supporting a contextual-shift rather than direct name-matching explanation.

F Query Format Robustness Check

To test whether the function-word TCD amplification pattern generalizes beyond our primary query format (“Who is an experienced {profession}?”), we evaluate five query formats with varying function/content word structures ($n = 700$ tests each, crossing 10 name pairs, 10 professions, and 7 templates):

1. **Q1 (Original):** “Who is an experienced {profession}?” (4 function words: who, is, an, ?)
2. **Q2 (Prepositional clause):** “[{profession} with extensive professional experience” (1 function word: with)

Format	#Func	#Cont	Ratio	d	Wilcoxon p
Q1 (Original)	4.0	2.2	$1.73 \times$	0.67	3.26×10^{-53}
Q2 (Prepos.)	1.0	4.2	$1.50 \times$	0.52	2.06×10^{-19}
Q3 (Imperative)	2.0	4.2	$1.55 \times$	0.61	6.02×10^{-33}
Q4 (Keyword)	0.0	4.2	—	0.00	1.00
Q5 (Question)	3.0	3.2	$1.58 \times$	0.66	2.69×10^{-45}

Table 9: Query format robustness check results ($n = 700$ per format). Across all syntax structures containing function words, the score-sensitivity amplification holds (effect size $d = 0.52$ – 0.67 ; Wilcoxon signed-rank $p < 10^{-18}$). These ratios are computed on the robustness subset (10 name pairs $\times$ 10 professions $\times$ 7 templates) and are therefore not directly comparable to the full 55,440-test controlled template aggregate ratio.

3. **Q3 (Imperative):** “Find a qualified {profession} for hire” (2 function words: a, for)
4. **Q4 (Keyword-only):** “experienced {profession} resume qualifications” (0 function words)
5. **Q5 (Question variant):** “What makes a good {profession}?” (3 function words: what, a, ?)

Table 9 reports the results. Across all query formats that contain function words (Q1, Q2, Q3, Q5), the Func/Cont TCD ratio is consistently and significantly greater than 1.0 (ranging from $1.50 \times$ to $1.73 \times$; Wilcoxon signed-rank $p < 10^{-18}$; Cohen’s $d = 0.52$ – 0.67). In the keyword-only baseline (Q4), the ratio is naturally N/A because no function words are present. This suggests that the function-word sensitivity pattern is robust across diverse syntactic structures, rather than an artifact of a single query grammar.

G Multiple Comparisons Correction

To evaluate the robustness of our statistical claims, we apply multiple-comparison corrections. First, we examine cell-by-cell hypothesis tests across a representative subset of 7 templates (yielding 231 individual profession $\times$ template tests, each evaluated over the 5 counterfactual name pairs). Because each individual cell has a small sample size ($n = 5$ paired observations), the non-parametric Wilcoxon signed-rank test is highly underpowered at this fine-grained level—the absolute minimum possible uncorrected p -value is $1/2^5 \approx 0.031$. Under this setup, 43 out of 231 individual cells are statistically significant before correction (where all 5 name pairs consistently exhibit higher function-word TCD). However, after applying multiple-comparison correction (Benjamini-Hochberg FDR,

Holm-Bonferroni, or strict Bonferroni), the number of significant individual cells drops to exactly zero. This correction is not intended as the primary evidence for or against the aggregate function-word effect; rather, it illustrates that cell-level hypothesis testing with $n = 5$ is underpowered and should not replace cluster-level inference. This fine-grained power limitation highlights why we emphasize aggregate-level analyses, cross-template bootstrap confidence intervals, and mixed-effects models (which leverage the full 55,440-observation scale) rather than relying on cell-by-cell significance testing.

Second, as a supplementary directional analysis, we evaluate the pairwise race-associated score differences across the 33 professions (evaluating whether specific racial categories systematically score higher than others). Before correction, 43 directional comparisons are statistically significant at $\alpha = 0.05$. However, after applying multiple-comparison correction (BH-FDR or Holm), the number of statistically significant pairwise differences drops to exactly zero.

This corrected directional analysis reinforces our caution against interpreting name-swap sensitivity as a global demographic ordering, and supports the distinction between absolute score sensitivity (TCD) and directional group harm. Name-swap sensitivity alone should not be conflated with directional societal disadvantage.

H Synthetic Candidate-Pool Ranking Simulation

To evaluate the practical ranking consequences of score sensitivity, we construct synthetic candidate pools with qualification variation. This setup is designed to approximate realistic qualification variance, but does not represent a live production retrieval system.

Pool Construction. For each of the 8 professions, we generate candidate pools of size $N \in \{10, 20, 50\}$. Candidate profiles are generated using 10 structurally distinct qualification templates (e.g., describing years of clinical experience, advanced credentials, or team leadership size) mixed with random values (e.g., years $\in [5, 25]$, university prestige, and number of published papers) to introduce realistic score variance. For each pool, we perform name swaps using our 5 race-associated name pairs (Emily $\leftrightarrow$ Lakisha, Brett $\leftrightarrow$ Jamal, Sarah

N	Any Δ	Mean	Top-1	Top-3	Top-10
10	36.6% [33.7,39.7]	0.58 [0.52,0.64]	8.5% [5.0,12.5]	9.8% [7.5,12.2]	N/A
20	55.0% [52.6,57.1]	1.21 [1.12,1.29]	21.0% [15.5,26.5]	7.8% [5.8,10.2]	10.4% [9.2,11.6]
50	75.1% [73.7,76.4]	3.07 [2.92,3.21]	41.0% [34.5,48.0]	18.3% [14.7,22.0]	7.9% [6.4,9.5]

Table 10: Complete ranking impact metrics across pool sizes under counterfactual identity swaps (200 trials per pool size with 95% bootstrap confidence intervals in brackets). Top- k set changes refer to membership changes at that cutoff; a Top-3 change need not imply a Top-10 change if the swapped candidates remain within the top-10.

$\leftrightarrow$ Tyrone, Allison $\leftrightarrow$ Tanisha, Neil $\leftrightarrow$ Leroy) to create identical counterfactual pools.

Displacement and Top- k Set Changes. Table 10 reports the complete set of rank displacement and top- k change rates. Across pool sizes of 10, 20, and 50 documents, identity swaps lead to average rank displacements of 0.58, 1.21, and 3.07 positions, respectively. The set of top candidates is also impacted: the Top-3 set changes in 9.8%, 7.8%, and 18.3% of cases for pools of 10, 20, and 50 candidates, while the Top-10 set changes in 10.4% of cases for pools of 20 candidates.

Score Margin Flip Analysis. A key concern in retrieval fairness is whether score sensitivity can trigger a rank flip between adjacent candidates. When candidates have similar qualifications, their initial score margin (difference in ColBERT score) is extremely small. In our simulations, when the score margin between candidates is small (< 0.05), even a minor score sensitivity shift (e.g., TCD ≈ 0.027) is mathematically sufficient to flip their relative rank. Table 11 reports the empirical adjacent margin flip statistics. As candidate density increases (pool size 50), a massive 83.9% of adjacent pairs fall into the < 0.05 margin range, where flips occur at a very high rate (25.6%–31.0%).

I Computational Details

All models are publicly available under permissive licenses (MIT/Apache 2.0). ColBERTv2 uses a BERT-base backbone (~ 110 M parameters). DPR (~ 110 M), Contriever (~ 110 M), SPLADE (~ 110 M), and the cross-encoder MiniLM-L-6-v2 (~ 22 M) are all run at inference only, with no fine-tuning or hyperparameter search.

All experiments were run on a single Apple M-series CPU (no GPU) in approximately 50 minutes total. The main ColBERT experiment (55,440 tests) takes ~ 22 minutes; each multi-model comparison

Pool Size	Margin Bin	# Pairs	% Pairs	Flip Rate
10	< 0.01	211	11.7%	25.1%
	$0.01 - 0.05$	502	27.9%	31.7%
	$0.05 - 0.10$	334	18.6%	20.4%
	> 0.10	753	41.8%	6.9%
20	< 0.01	906	23.8%	29.9%
	$0.01 - 0.05$	1391	36.6%	29.8%
	$0.05 - 0.10$	645	17.0%	24.2%
	> 0.10	858	22.6%	7.0%
50	< 0.01	4549	46.4%	25.6%
	$0.01 - 0.05$	3672	37.5%	31.0%
	$0.05 - 0.10$	868	9.9%	23.2%
	> 0.10	711	7.3%	8.0%

Table 11: Adjacent score margin flip statistics across candidate pool sizes, showing pair distribution and actual rank flip rates per margin bin.




Figure 8: **Aggregate identity sensitivity across architectures.** Comparison of aggregate score sensitivity (SS) across different neural retrieval architectures ($n = 1,155$ per model). While all models exhibit non-zero aggregate sensitivity, only ColBERT supports exact token-level decomposition.

(1,155 tests) takes ~ 5 minutes; the naturalistic-text validation (8,000 tests) takes ~ 7 minutes. The 200-trial expanded synthetic ranking simulations and query-ablation probes add approximately 6 minutes of runtime. Runtime excludes optional plotting and bootstrap resampling.